

Project Objective

Develop machine learning models to predict molecular toxicity across 12 biological assays using SMILES representations, enabling in-ico screening for drug discovery and chemical safety assessment.

Dataset Overview

- **Source:** Tox21 Challenge Dataset
- **Size:** 7,831 molecules
- **Features:** SMILES strings
- **Targets:** 12 binary toxicity assays (NR and SR pathways)
- **Challenge:** High class imbalance (2-16% positive class) and missing labels (7-26%)

The workflow includes:

- Loading and inspecting the Tox21 dataset
- Handling missing assay values and class imbalance
- Generating molecular descriptors and fingerprints from SMILES using RDKit
- Exploring descriptor distributions before and after normalization
- Training multiple machine learning models for each toxicity assay
- Selecting the best model per assay based on cross-validated F1 score

The final outcome is a set of optimized binary classification models that can predict toxic or non-toxic behavior of molecules for individual biological targets.

In [102...]

```
# -- General --
import joblib
import os

# -- EDA and Preprocessing --
import pandas as pd
import numpy as np
```

```
# -- Visualization --
import seaborn as sns
import matplotlib.pyplot as plt

# -- RDKit --
from rdkit import Chem
from rdkit.Chem import Draw, rdMolDescriptors, Lipinski, Descriptors, Crippen
from rdkit import DataStructs
from rdkit.Chem.rdBFPGenerator import GetMorganGenerator

# -- Machine Learning --
from sklearn.model_selection import train_test_split
from sklearn.base import clone
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.linear_model import LogisticRegression
from lightgbm import LGBMClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.pipeline import Pipeline

# -- Preprocessing --
from sklearn.preprocessing import PowerTransformer, StandardScaler

# -- Metrics --
from sklearn.metrics import f1_score, confusion_matrix, classification_report

import warnings

warnings.filterwarnings(
    "ignore",
    message="X does not have valid feature names",
    category=UserWarning
)
```

1. Loading Data

In [103...]

```
# Loading Data
df_tox = pd.read_csv('../data/tox21.csv')
```

In [104...]

```
print(f'The rows in the dataset are {df_tox.shape[0]} and the columns are {df_tox.shape[1]}!\\n')
print(f'The columns in the dataset are:')

for no, col in enumerate(df_tox.columns, start=1):
    print(f'{no}. {col}')
```

The rows in the dataset are 7831 and the columns are 14!

The columns in the dataset are:

1. NR-AR
2. NR-AR-LBD
3. NR-AhR
4. NR-Aromatase
5. NR-ER
6. NR-ER-LBD
7. NR-PPAR-gamma
8. SR-ARE
9. SR-ATAD5
10. SR-HSE
11. SR-MMP
12. SR-p53
13. mol_id
14. smiles

Tox21 Dataset Column Definitions

Nuclear Receptor (NR) Pathways

- **NR-AR** : Tests if a chemical affects the Androgen (male hormone) receptor.
- **NR-AR-LBD** : Tests if a chemical binds to the active site of the Androgen receptor.
- **NR-AhR** : Tests if a chemical acts like toxic environmental pollutants.
- **NR-Aromatase** : Tests if a chemical blocks estrogen formation.
- **NR-ER** : Tests if a chemical affects the Estrogen (female hormone) receptor.
- **NR-ER-LBD** : Tests if a chemical binds to the active site of the Estrogen receptor.

- **NR-PPAR-gamma** : Tests if a chemical affects fat and metabolism control.

Stress Response (SR) Pathways

- **SR-ARE** : Tests if a chemical causes oxidative (cell) stress.
- **SR-ATAD5** : Tests if a chemical damages DNA during replication.
- **SR-HSE** : Tests if a chemical causes protein stress or misfolding.
- **SR-MMP** : Tests if a chemical damages mitochondria (cell energy).
- **SR-p53** : Tests if a chemical triggers DNA damage response or cell death.

🔍 Key Observations

- The dataset contains 7,831 molecules, each identified by a unique `mol_id` and a corresponding `smiles` string.
- There are 12 toxicity assay columns, covering Nuclear Receptor (NR) and Stress Response (SR) pathways.
- The dataset structure is suitable for multi-target binary classification, where each assay is treated as an independent prediction task.

2. Data Inspection

```
In [105...]: df_tox.shape
```

```
Out[105...]: (7831, 14)
```

```
In [106...]: df_tox.head()
```

Out[106...]

	NR-AR	NR-AR-LBD	NR-AhR	NR-Aromatase	NR-ER	NR-ER-LBD	NR-PPAR-gamma	SR-ARE	SR-ATAD5	SR-HSE	SR-MMP	SR-p53	mol_id	
0	0.0	0.0	1.0	NaN	NaN	0.0	0.0	1.0	0.0	0.0	0.0	0.0	TOX3021	CCOc1ccc2nc(=O)Nc1c(C)c2ccccc2C=C2CCCC4=CCCC[C@H]2C=C4
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	NaN	0.0	NaN	0.0	0.0	TOX3020	CCN1C(=O)NC(c2cccc2)C1=O
2	NaN	NaN	NaN	NaN	NaN	NaN	NaN	0.0	NaN	0.0	NaN	NaN	TOX3024	CC[C@]1(O)CC[C@H]2[C@@H]3CCCC4=CCCC[C@H]2C=C4C3
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	NaN	0.0	NaN	0.0	0.0	TOX3027	CCCN(CC)C(CC)C(=O)Nc1c(C)c2ccccc2C=C2CCCC4=CCCC[C@H]2C=C4
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	TOX20800	CC(O)(P(=O)(O)O)P(=O)(O)O

◀ ▶

In [107...]

```
# Reordering columns
cols = df_tox.columns.tolist()
cols = ['mol_id', 'smiles'] + [c for c in cols if c != 'smiles' and c != 'mol_id']
df_tox = df_tox[cols]
```

In [108...]

```
df_tox.head()
```

Out[108...]

	mol_id	smiles	NR-AR	NR-AR-LBD	NR-AhR	NR-Aromatase	NR-ER	NR-ER-LBD	NR-PPAR-gamma	SR-ARE	SR-ATAD5	SR-HSE	SR-MMP	SR-p53	
0	TOX3021	CCOc1ccc2nc(=O)=Osc2c1	0.0	0.0	1.0	NaN	NaN	0.0	0.0	1.0	0.0	0.0	0.0	0.0	
1	TOX3020	CCN1C(=O)NC(c2cccc2)C1=O	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	N
2	TOX3024	CC[C@]1(O)CC[C@H]2[C@@H]3CCCC4=CCCC[C@H]2C=C4C3	NaN	NaN	NaN	NaN	NaN	NaN	NaN	0.0	0.0	0.0	0.0	0.0	
3	TOX3027	CCCN(CC)C(CC)C(=O)Nc1c(C)c2ccccc2C=C2CCCC4=CCCC[C@H]2C=C4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	N
4	TOX20800	CC(O)(P(=O)(O)O)P(=O)(O)O	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	

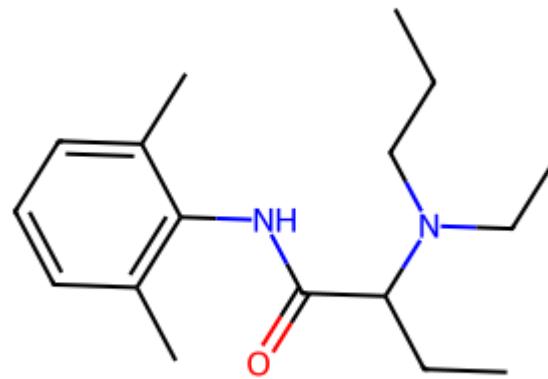
◀ ▶

In [109...]

```
# Example of Visualization of a molecule from DataSet
mol = Chem.MolFromSmiles("CCCN(CC)C(CC)C(=O)Nc1c(C)c2ccccc2C=C2CCCC4=CCCC[C@H]2C=C4")
```

```
Draw.MolToImage(mol)
```

Out[109...]



In [110...]

```
df_tox.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7831 entries, 0 to 7830
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   mol_id       7831 non-null   object  
 1   smiles       7831 non-null   object  
 2   NR-AR        7265 non-null   float64 
 3   NR-AR-LBD   6758 non-null   float64 
 4   NR-AhR      6549 non-null   float64 
 5   NR-Aromatase 5821 non-null   float64 
 6   NR-ER        6193 non-null   float64 
 7   NR-ER-LBD   6955 non-null   float64 
 8   NR-PPAR-gamma 6450 non-null   float64 
 9   SR-ARE       5832 non-null   float64 
 10  SR-ATAD5    7072 non-null   float64 
 11  SR-HSE       6467 non-null   float64 
 12  SR-MMP      5810 non-null   float64 
 13  SR-p53      6774 non-null   float64 
dtypes: float64(12), object(2)
memory usage: 856.6+ KB
```

In [111... df_tox.columns

```
Out[111... Index(['mol_id', 'smiles', 'NR-AR', 'NR-AR-LBD', 'NR-AhR', 'NR-Aromatase',
 'NR-ER', 'NR-ER-LBD', 'NR-PPAR-gamma', 'SR-ARE', 'SR-ATAD5', 'SR-HSE',
 'SR-MMP', 'SR-p53'],
 dtype='object')
```

```
In [112... target_columns = ['NR-AR', 'NR-AR-LBD', 'NR-AhR', 'NR-Aromatase',
 'NR-ER', 'NR-ER-LBD', 'NR-PPAR-gamma', 'SR-ARE',
 'SR-ATAD5', 'SR-HSE', 'SR-MMP', 'SR-p53']

summary = pd.DataFrame({
    "null_count": df_tox[target_columns].isna().sum(),
    "non_null_count": df_tox[target_columns].notna().sum(),
    "null_%": (df_tox[target_columns].isna().mean() * 100).round(2),

    # class counts (ignore NaNs automatically)
    "count_0": (df_tox[target_columns] == 0).sum(),
    "count_1": (df_tox[target_columns] == 1).sum(),
```

```

# percentage of positive class
"%_ones": ((df_tox[target_columns] == 1).sum()
            / df_tox[target_columns].notna().sum()
            * 100).round(2)

}).sort_values("null_count", ascending=False)

summary

```

Out[112...]

	null_count	non_null_count	null_%	count_0	count_1	%_ones
SR-MMP	2021	5810	25.81	4892	918	15.80
NR-Aromatase	2010	5821	25.67	5521	300	5.15
SR-ARE	1999	5832	25.53	4890	942	16.15
NR-ER	1638	6193	20.92	5400	793	12.80
NR-PPAR-gamma	1381	6450	17.64	6264	186	2.88
SR-HSE	1364	6467	17.42	6095	372	5.75
NR-AhR	1282	6549	16.37	5781	768	11.73
NR-AR-LBD	1073	6758	13.70	6521	237	3.51
SR-p53	1057	6774	13.50	6351	423	6.24
NR-ER-LBD	876	6955	11.19	6605	350	5.03
SR-ATAD5	759	7072	9.69	6808	264	3.73
NR-AR	566	7265	7.23	6956	309	4.25

In [113...]

df_tox.columns

```

Out[113...]: Index(['mol_id', 'smiles', 'NR-AR', 'NR-AR-LBD', 'NR-AhR', 'NR-Aromatase',
       'NR-ER', 'NR-ER-LBD', 'NR-PPAR-gamma', 'SR-ARE', 'SR-ATAD5', 'SR-HSE',
       'SR-MMP', 'SR-p53'],
      dtype='object')

```

```
In [114... df_tox['smiles'].duplicated().sum()
```

```
Out[114... np.int64(0)
```

🔍 Key Observations

- Each assay is a binary label where:
 - 0 indicates non-toxic (inactive)
 - 1 indicates toxic (active)
- All assay columns contain missing values, indicating that not every molecule was tested against every biological target.
- The percentage of missing values varies across assays, with some targets having more than 25% missing data.
- The class distribution is highly imbalanced:
 - Most molecules are labeled as non-toxic (0)
 - Toxic (1) samples are comparatively rare
- This imbalance highlights the need for appropriate evaluation metrics (F1 score) and balanced learning strategies during modeling.

3. Feature Engineering

```
In [115... def calculate_descriptors_safe(smiles):  
    """  
        Calculates multiple molecular descriptors from a SMILES string.  
        Returns a pandas Series so it can add multiple columns at once.  
        Output schema is IDENTICAL to the original function.  
    """  
    morgan_gen = GetMorganGenerator(radius=3, fpSize=2048)  
    try:  
        # 1. Safely create and sanitize molecule  
        mol = Chem.MolFromSmiles(smiles, sanitize=True)  
        if mol is None:  
            raise ValueError("Invalid SMILES")  
  
        # 2. Remove explicit hydrogens (fixes H-without-neighbors warning)  
        mol = Chem.RemoveHs(mol)
```

```
# 3. Calculate descriptors (same as before)
return pd.Series({
    'tpsa': rdMolDescriptors.CalcTPSA(mol),
    'mw': Descriptors.MolWt(mol),
    'hbd': Descriptors.NumHDonors(mol),
    'hba': Descriptors.NumHAcceptors(mol),
    'rot': Descriptors.NumRotatableBonds(mol),
    'logP': Crippen.MolLogP(mol),
    'mr': Crippen.MolMR(mol),
    'hac': Lipinski.HeavyAtomCount(mol),
    'fp': morgan_gen.GetFingerprint(mol), # 🔥 NO deprecation warning
    'ring_count': rdMolDescriptors.CalcNumRings(mol)
})

except Exception:
    # 4. Fail safely (same NaN structure as original)
    return pd.Series({
        'tpsa': np.nan,
        'mw': np.nan,
        'hbd': np.nan,
        'hba': np.nan,
        'rot': np.nan,
        'logP': np.nan,
        'mr': np.nan,
        'hac': np.nan,
        'fp': np.nan,
        'ring_count': np.nan
    })

# Apply the function (Creating all columns in one go)
# This returns a new DataFrame with the descriptor columns
descriptor_cols = df_tox['smiles'].apply(calculate_descriptors_safe)

# Join these new columns back to your original DataFrame
df_tox = pd.concat([df_tox, descriptor_cols], axis=1)

# Check for failures (just checking one column is enough since they all fail together)
print(f"Failed rows: {df_tox['tpsa'].isna().sum()}")
```

```
[14:00:42] WARNING: not removing hydrogen atom without neighbors
[14:00:42] WARNING: not removing hydrogen atom without neighbors
[14:00:47] Explicit valence for atom # 8 Al, 6, is greater than permitted
[14:00:49] Explicit valence for atom # 3 Al, 6, is greater than permitted
[14:00:49] Explicit valence for atom # 4 Al, 6, is greater than permitted
[14:00:52] Explicit valence for atom # 4 Al, 6, is greater than permitted
[14:00:53] Explicit valence for atom # 9 Al, 6, is greater than permitted
[14:00:53] Explicit valence for atom # 5 Al, 6, is greater than permitted
[14:00:56] Explicit valence for atom # 16 Al, 6, is greater than permitted
[14:00:58] Explicit valence for atom # 20 Al, 6, is greater than permitted
Failed rows: 8
```

```
In [116...]: # Filter for rows where TPSA is NaN (empty)
invalid_rows = df_tox[df_tox['hbd'].isna()].index
```

```
# Display just the SMILES strings that failed
print("Invalid SMILES strings:")
print(invalid_rows)
```

```
df_tox = df_tox.drop(index=invalid_rows).reset_index(drop=True)
```

```
Invalid SMILES strings:
```

```
Index([1322, 2290, 2297, 3558, 4565, 4649, 5538, 6723], dtype='int64')
```

```
In [117...]: # Filter for rows where TPSA is NaN (empty)
invalid_rows = df_tox[df_tox['hbd'].isna()].index
```

```
# Display just the SMILES strings that failed
print("Invalid SMILES strings:")
print(invalid_rows)
```

```
Invalid SMILES strings:
```

```
Index([], dtype='int64')
```

```
In [118...]: df_tox.shape
```

```
Out[118...]: (7823, 24)
```

```
In [119...]: df_tox.head(1)
```

Out[119...]

	mol_id	smiles	NR-AR	NR-AR-LBD	NR-AhR	NR-Aromatase	NR-ER	NR-ER-LBD	NR-PPAR-gamma	SR-ARE	...	tpsa	mw	hbd	hba	rot	logP	mr
0	TOX3021	CCOc1ccc2nc(S(N)(=O)=O)sc2c1	0.0	0.0	1.0	NaN	NaN	0.0	0.0	1.0	...	82.28	258.324	1.0	5.0	3.0	1.3424	62.1622

1 rows × 24 columns



Why these descriptors?

- **TPSA, LogP, MW:** Drug-likeness indicators (Lipinski's Rule of Five)
- **HBD/HBA:** Membrane permeability predictors
- **Morgan Fingerprints:** Capture structural similarity and substructure presence
- **Rotatable Bonds:** Molecular flexibility indicator

🔍 Key Observations

- SMILES strings were converted into numerical features because machine learning models cannot directly interpret chemical strings.
- Molecular descriptors such as TPSA, molecular weight, hydrogen bond counts, and rotatable bonds capture physicochemical properties of molecules.

- Morgan fingerprints (2048 bits) encode structural sub-patterns and substructures present in each molecule.
- A small number of invalid SMILES failed RDKit processing and were safely removed to maintain data quality.
- The combination of scalar descriptors and fingerprints provides both global chemical properties and detailed structural information for toxicity prediction.

4. Descriptor Distribution Analysis

```
In [120]: cols = ['tpsa', 'mw', 'logp', 'hbd', 'hba', 'rot', 'mr', 'hac', 'ring_count']
```

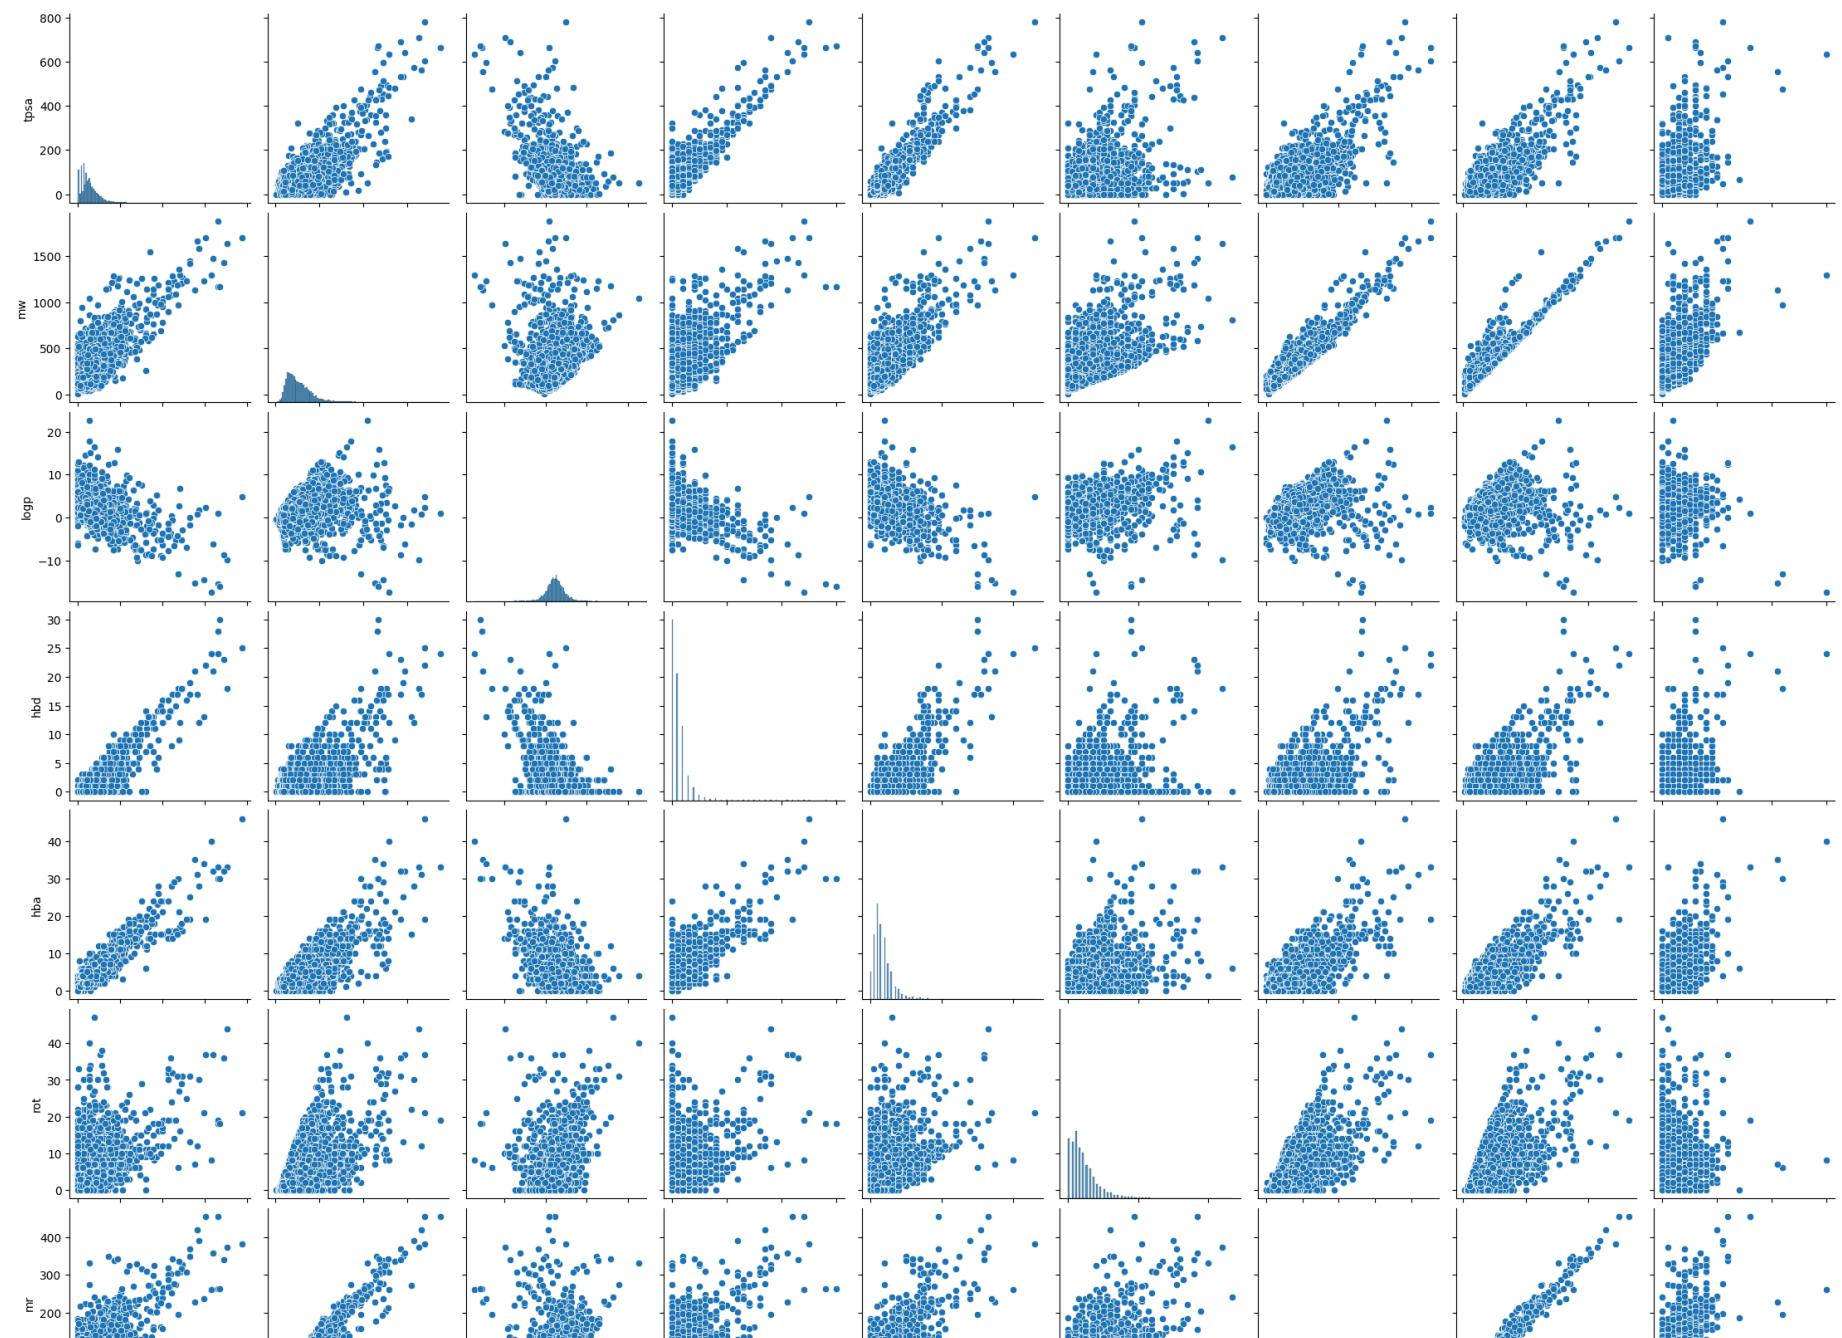
```
In [121]: sns.pairplot(df_tox[cols])

# Add title
plt.suptitle('Pairwise Relationship between Descriptors', y=1.02)

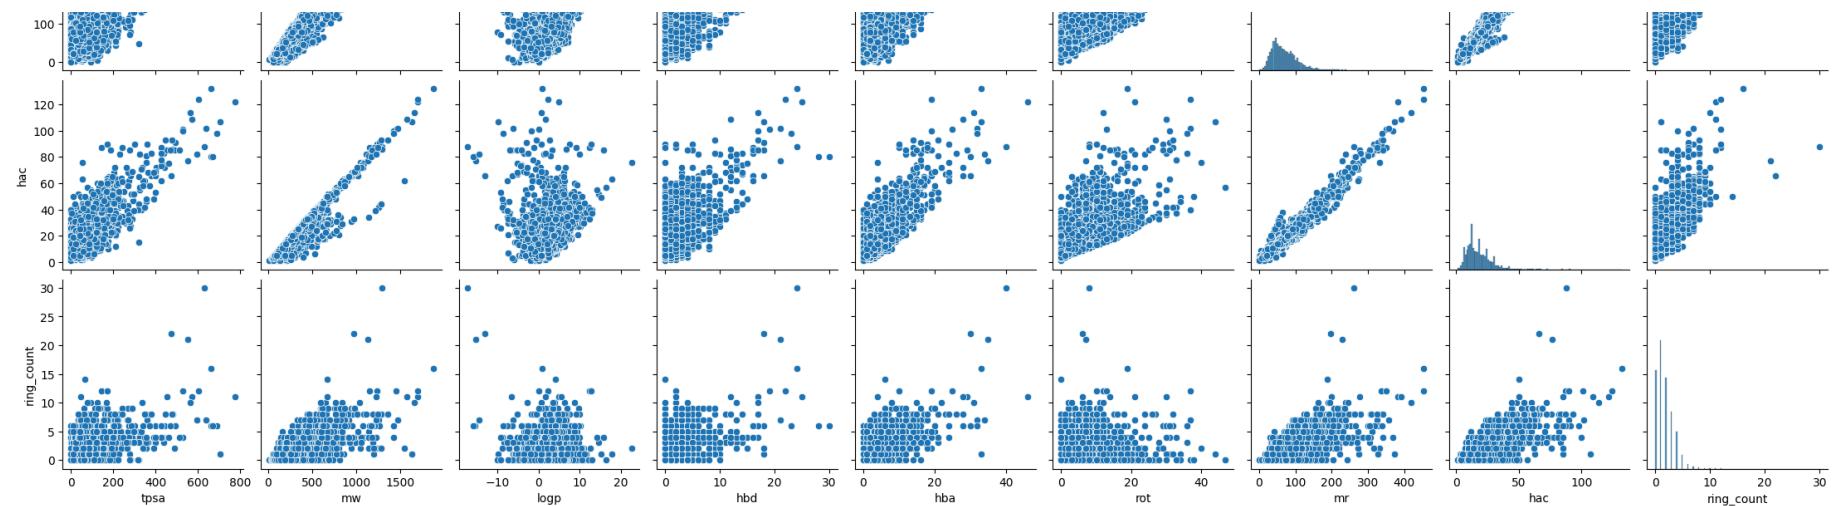
# Show plot
plt.show()
```

1.EDA

Pairwise Relationship between Descriptors



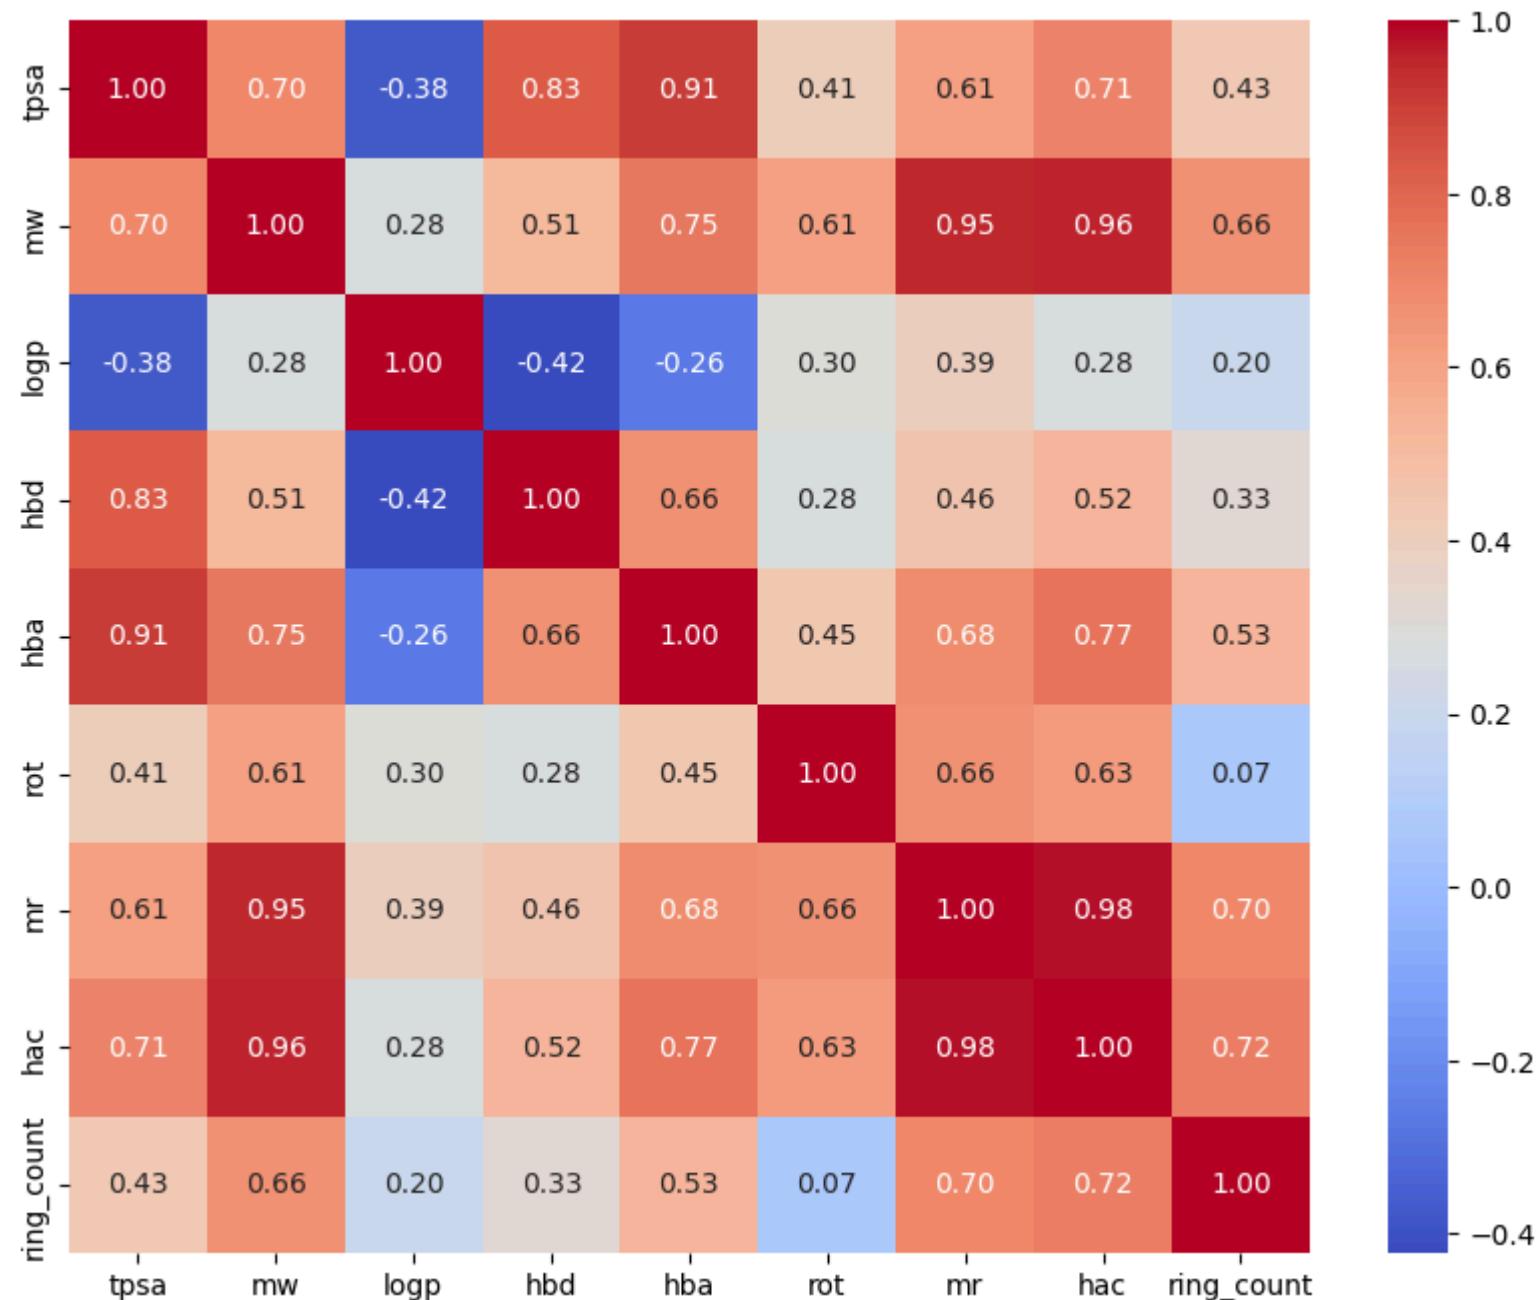
1.EDA



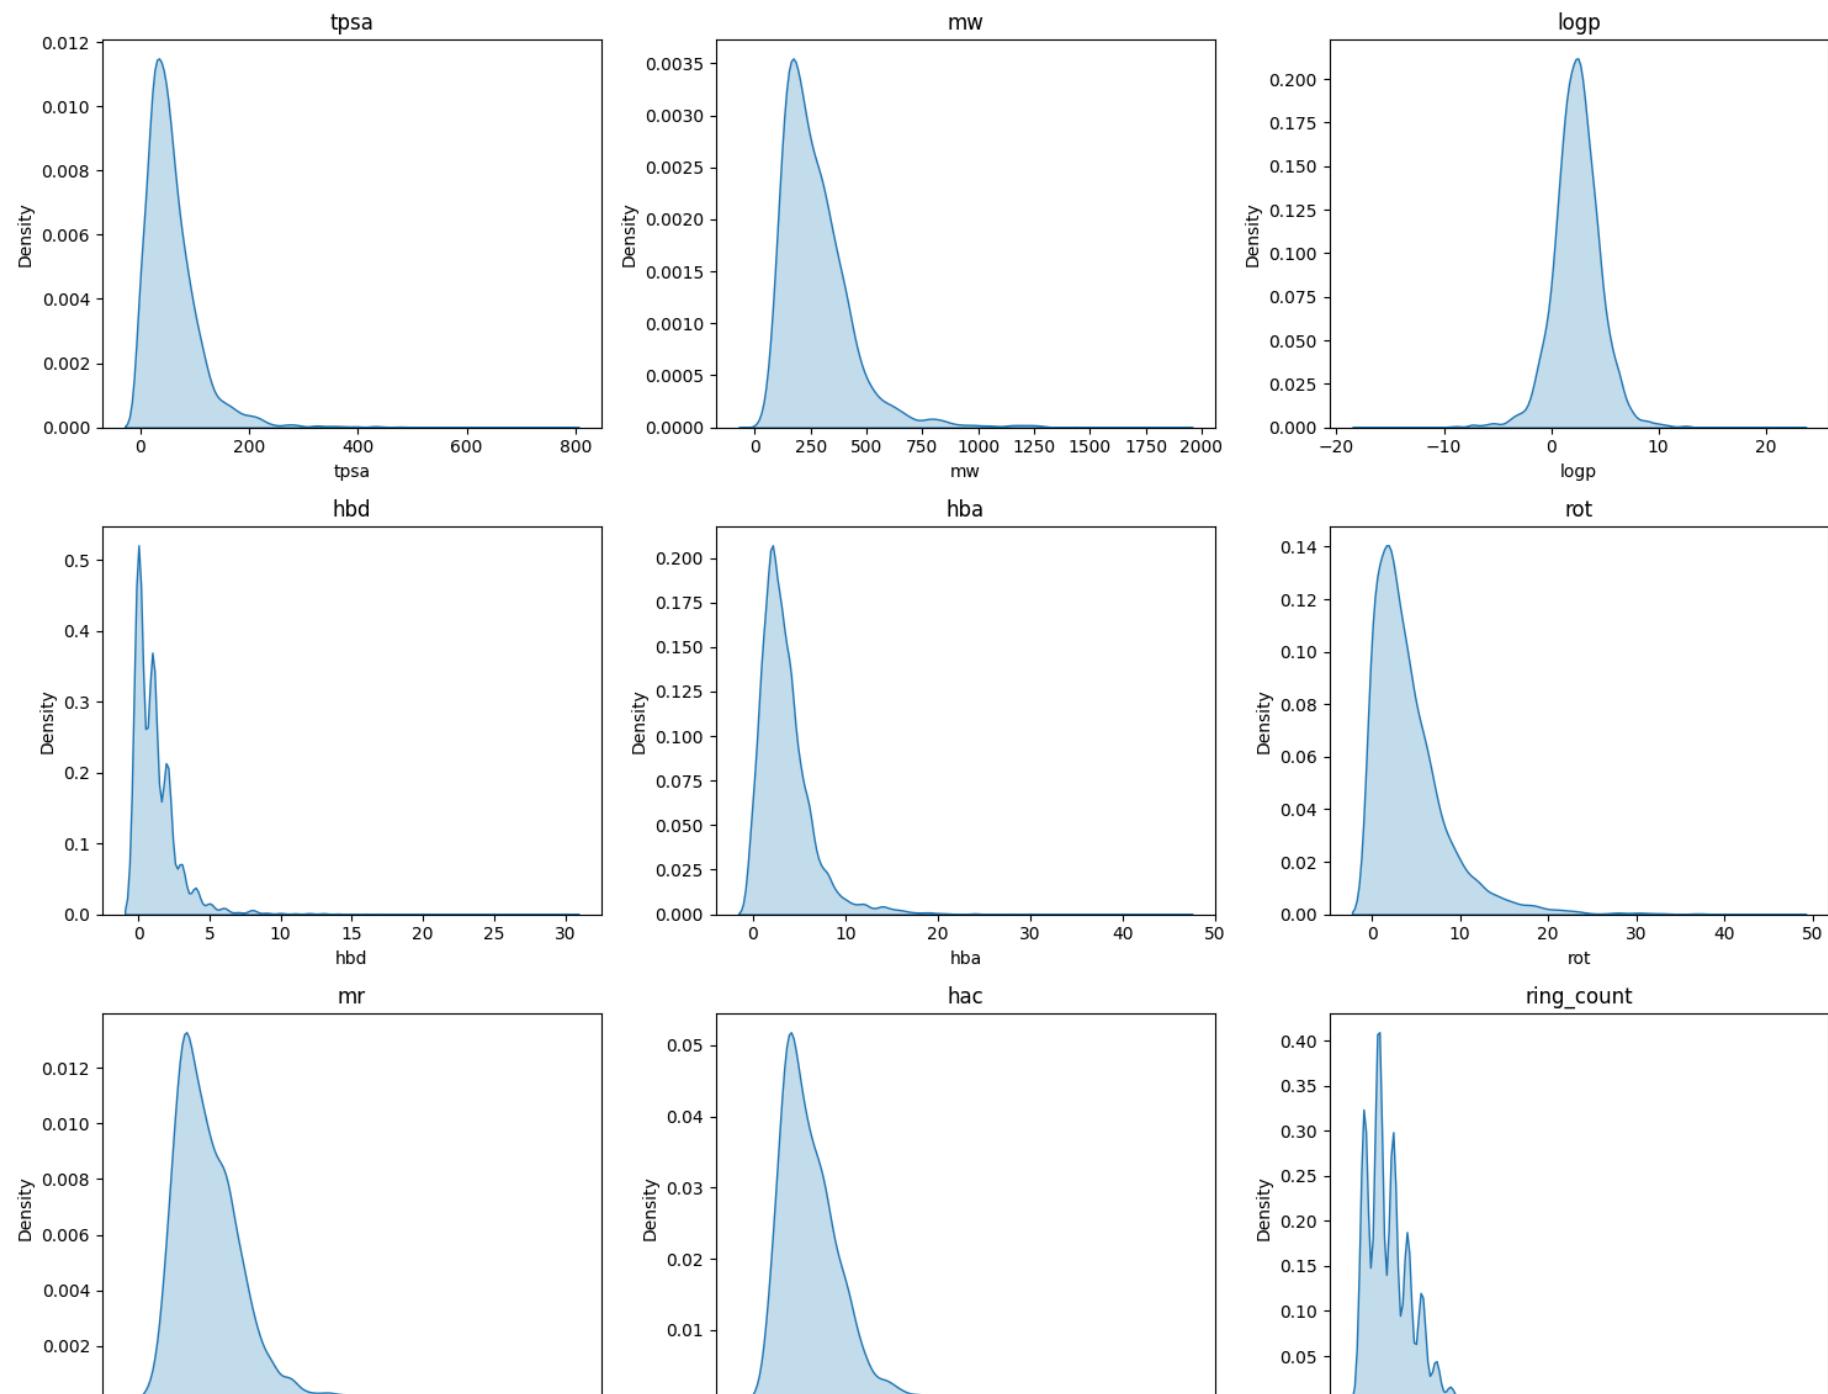
In [122]:

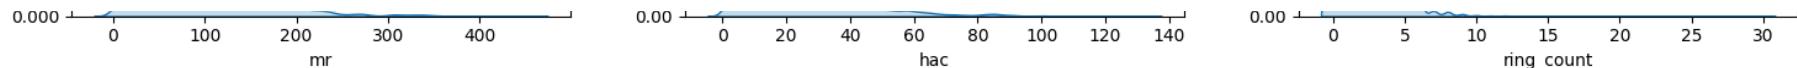
```
corr_matrix = df_tox[cols].corr()

plt.figure(figsize=(10,8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.show()
```



```
In [123...]:  
fig, axes = plt.subplots(3, 3, figsize=(15, 12))  
axes = axes.flatten()  
  
for i, col in enumerate(cols):  
    sns.kdeplot(data=df_tox, x=col, fill=True, ax=axes[i])  
    axes[i].set_title(col)  
  
# Remove empty subplots if cols < 9  
for j in range(len(cols), 9):  
    fig.delaxes(axes[j])  
  
plt.tight_layout()  
# plt.savefig(r'..\plots\1.Descriptor_without_normalization.png')  
plt.show()
```





🔍 Key Observations

- Most molecular descriptors show right-skewed distributions, which is common for chemical property data.
- Properties such as molecular weight, TPSA, and rotatable bonds span a wide range of values.
- The presence of skewness suggests that direct use of raw values may negatively impact certain machine learning models.
- Visual inspection confirms the need for normalization to stabilize variance and improve model learning.

5. Power Transformation & Normalization

In [124...]

```
n_rows = (len(cols) + 1) // 2
n_cols = 2

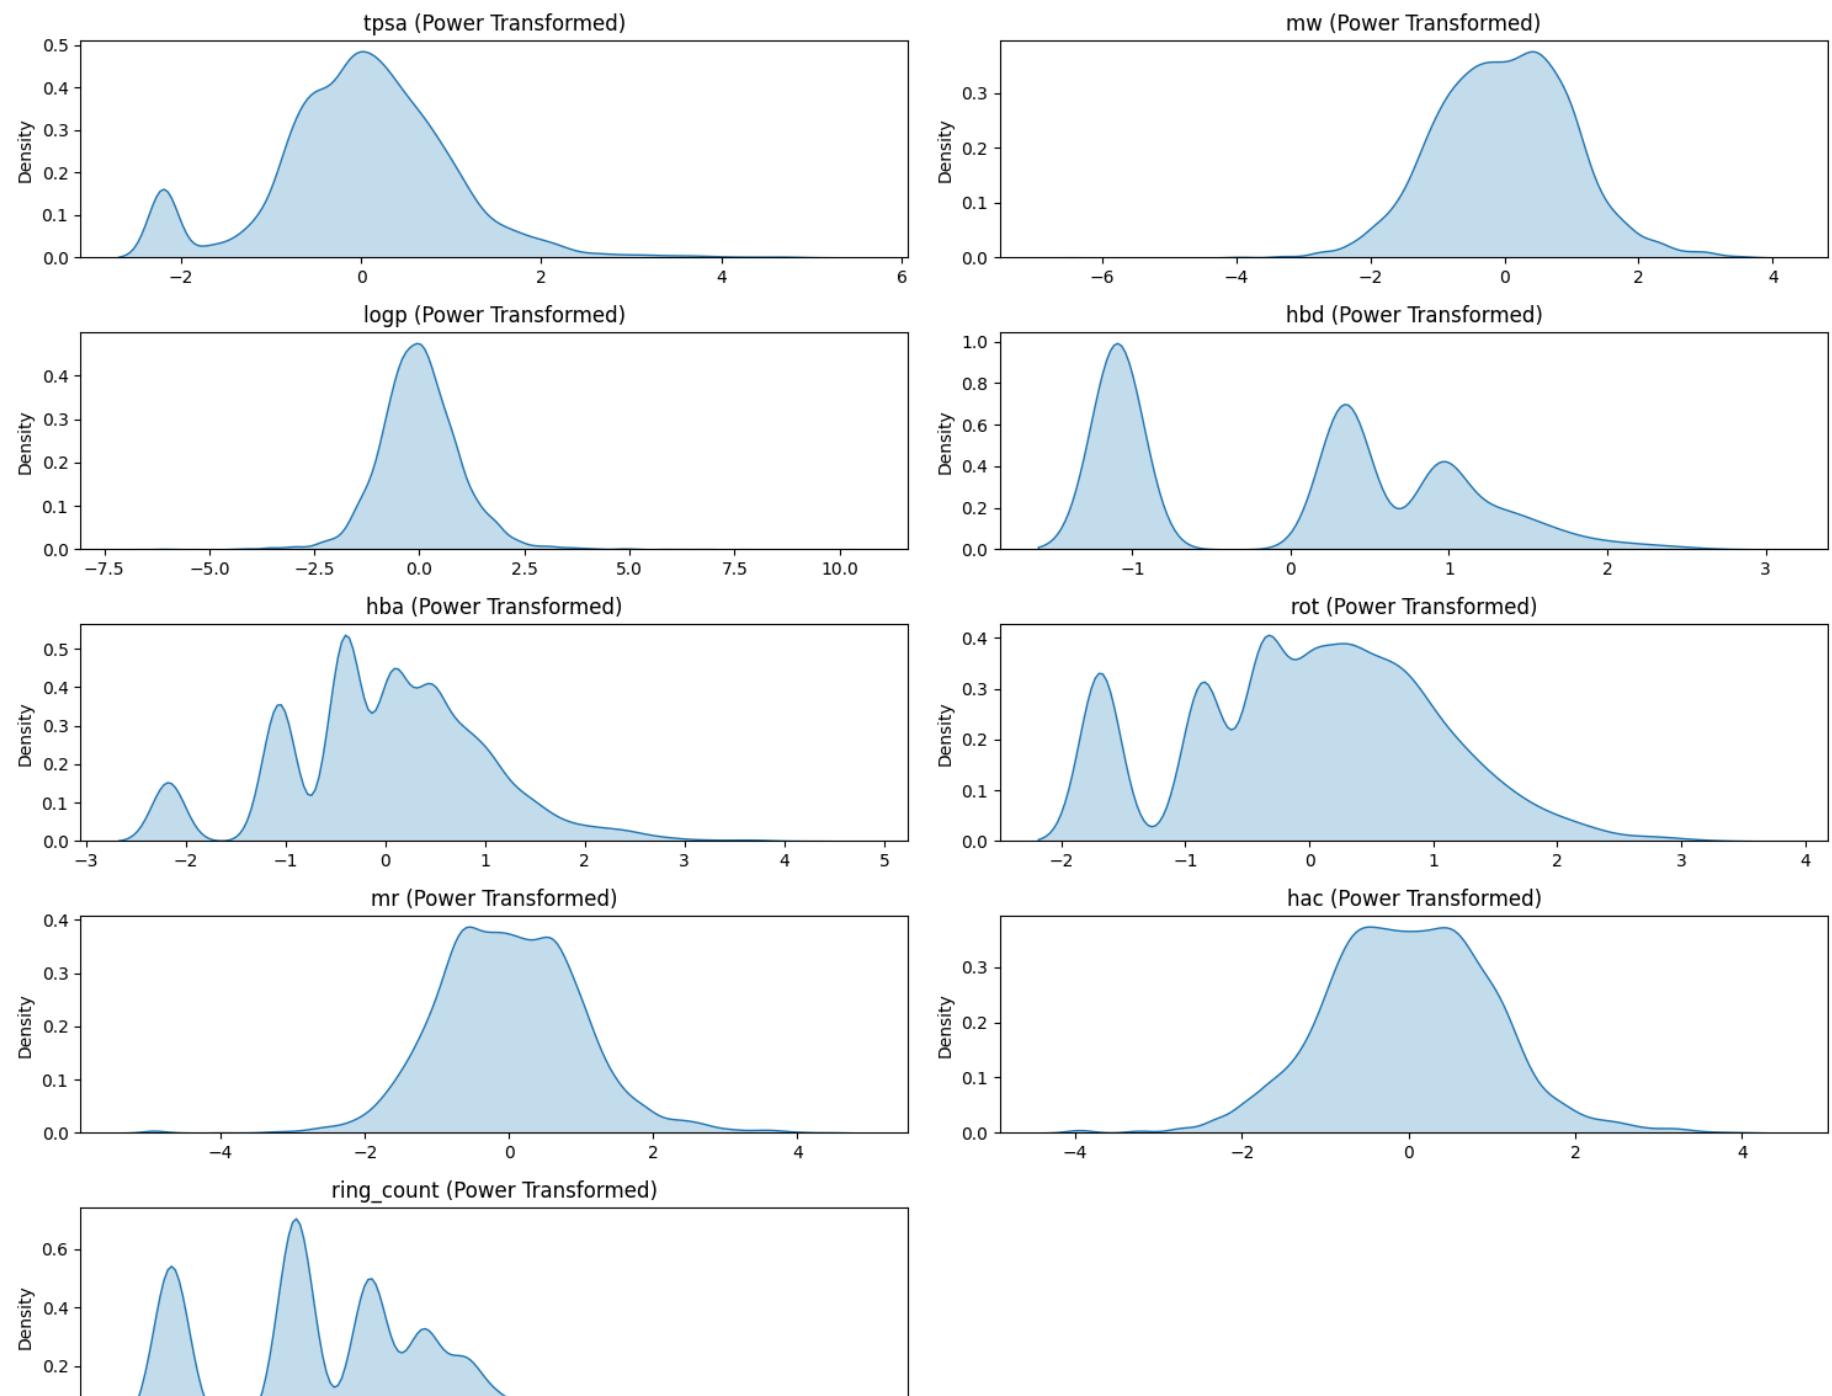
fig, axes = plt.subplots(n_rows, n_cols, figsize=(15, 12))
axes = axes.flatten()

for i, col in enumerate(cols):
    pt = PowerTransformer()
    transformed = pt.fit_transform(df_tox[[col]]).ravel()

    sns.kdeplot(x=transformed, fill=True, ax=axes[i])
    axes[i].set_title(f'{col} (Power Transformed)')

# remove unused axes
for j in range(len(cols), len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
# plt.savefig(r'..\plots\2.Descriptor_with_normalization.png')
plt.show()
```





In [125...]

```
# Converting fingerprint to numpy array
def to_array(fp):
    arr = np.zeros(fp.GetNumBits(), dtype=np.uint8)
    DataStructs.ConvertToNumpyArray(fp, arr)
    return arr

fp_array = df_tox['fp'].apply(to_array)

df_tox['fp'] = fp_array
```

🔍 Key Observations

- PowerTransformer (Yeo-Johnson) normalization was applied to selected molecular descriptors.
- After transformation, feature distributions become more symmetric and closer to a normal shape.
- Normalization is especially beneficial for distance-based and linear models such as Logistic Regression and KNN.
- Both normalized and non-normalized feature sets were evaluated to compare model performance.

6. Model Training & Evaluation

In [126...]

```
# Tox21 binary classification targets
targets = ['NR-AR', 'NR-AR-LBD', 'NR-AhR', 'NR-Aromatase',
           'NR-ER', 'NR-ER-LBD', 'NR-PPAR-gamma',
           'SR-ARE', 'SR-ATAD5', 'SR-HSE', 'SR-MMP', 'SR-p53'
          ]

# Scalar molecular descriptors
scalar_cols = ['tpsa', 'mw', 'hbd', 'hba', 'rot',
                'logp', 'ring_count', '#mr', 'hac'
               ]
```

In [127...]

```
models_config = {
    "LR": Pipeline([("scaler", StandardScaler()), ("clf", LogisticRegression(max_iter=1000))]),
```

```
"KNN": Pipeline([("scaler", StandardScaler()), ("clf", KNeighborsClassifier(n_neighbors=5))]),
"RF": RandomForestClassifier(random_state=24),
"XGB": XGBClassifier(random_state=24),
"LBGM_Bal": LBGMClassifier(class_weight="balanced", verbosity=-1, random_state=24),
"LBGM_Unbal": LBGMClassifier(is_unbalance=True, verbosity=-1, random_state=24)
}
```

```
In [128...]: MODEL_DIR = ".../models"
os.makedirs(MODEL_DIR, exist_ok=True)
```

```
In [129...]: results = {t: {} for t in targets}
best_models_store_normalize = {}

for target in targets:
    print(f"\n===== Target: {target} =====")

    df_sub = df_tox[df_tox[target].notna()]
    y = df_sub[target].values

    X_scalar = df_sub[scalar_cols].values
    X_fp = np.stack(df_sub["fp"].values)

    # ----- Train / Test Split -----
    Xs_train, Xs_test, Xfp_train, Xfp_test, y_train, y_test = train_test_split(
        X_scalar,
        X_fp,
        y,
        test_size=0.2,
        stratify=y,
        random_state=24
    )

    # ----- Power Transform (FIT ONLY ON TRAIN) -----
    pt = PowerTransformer(method="yeo-johnson")
    Xs_train_pt = pt.fit_transform(Xs_train)
    Xs_test_pt = pt.transform(Xs_test)

    # ----- Combine Scalar + FP -----
    X_train = np.hstack([Xs_train_pt, Xfp_train])
```

```
X_test = np.hstack([Xs_test_pt, Xfp_test])

best_f1 = -1
best_model = None
best_model_name = None

for name, model in models_config.items():
    print(f"\n>>> Model: {name}")

    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    # ----- Metrics -----
    report = classification_report(y_test, y_pred, digits=4)
    cm = confusion_matrix(y_test, y_pred)
    f1_class1 = f1_score(y_test, y_pred, pos_label=1)

    print("Classification Report:")
    print(report)

    print("Confusion Matrix:")
    print(cm)

    print(f"F1 (Class 1): {f1_class1:.4f}")

    results[target][name] = f1_class1

    if f1_class1 > best_f1:
        best_f1 = f1_class1
        best_model = model
        best_model_name = name

# ----- Save Best Model + Transformer -----
model_path = f"{MODEL_DIR}/{target}_best_model.joblib"
transformer_path = f"{MODEL_DIR}/{target}_power_transformer.joblib"

joblib.dump(best_model, model_path)
joblib.dump(pt, transformer_path)

best_models_store_normalize[target] = {
    "model_name": best_model_name,
```

```
    "model_path": model_path,  
    "transformer_path": transformer_path,  
    "best_f1_class1": best_f1  
}  
  
print(f"\n💾 Saved best model: {best_model_name}")  
print(f"🏆 Best F1 (Class 1): {best_f1:.4f}")
```

===== Target: NR-AR =====

>>> Model: LR

Classification Report:

	precision	recall	f1-score	support
0.0	0.9714	0.9791	0.9753	1390
1.0	0.4314	0.3548	0.3894	62
accuracy			0.9525	1452
macro avg	0.7014	0.6670	0.6823	1452
weighted avg	0.9484	0.9525	0.9503	1452

Confusion Matrix:

[[1361 29]

[40 22]]

F1 (Class 1): 0.3894

>>> Model: KNN

Classification Report:

	precision	recall	f1-score	support
0.0	0.9686	1.0000	0.9841	1390
1.0	1.0000	0.2742	0.4304	62
accuracy			0.9690	1452
macro avg	0.9843	0.6371	0.7072	1452
weighted avg	0.9700	0.9690	0.9604	1452

Confusion Matrix:

[[1390 0]

[45 17]]

F1 (Class 1): 0.4304

>>> Model: RF

Classification Report:

	precision	recall	f1-score	support
0.0	0.9726	0.9971	0.9847	1390
1.0	0.8519	0.3710	0.5169	62

accuracy			0.9704	1452
macro avg	0.9122	0.6840	0.7508	1452
weighted avg	0.9675	0.9704	0.9647	1452

Confusion Matrix:

```
[[1386  4]
 [ 39  23]]
```

F1 (Class 1): 0.5169

>>> Model: XGB

Classification Report:

	precision	recall	f1-score	support
0.0	0.9726	0.9950	0.9836	1390
1.0	0.7667	0.3710	0.5000	62
accuracy			0.9683	1452
macro avg	0.8696	0.6830	0.7418	1452
weighted avg	0.9638	0.9683	0.9630	1452

Confusion Matrix:

```
[[1383  7]
 [ 39  23]]
```

F1 (Class 1): 0.5000

>>> Model: LGBM_Bal

Classification Report:

	precision	recall	f1-score	support
0.0	0.9730	0.9856	0.9793	1390
1.0	0.5455	0.3871	0.4528	62
accuracy			0.9601	1452
macro avg	0.7592	0.6864	0.7161	1452
weighted avg	0.9548	0.9601	0.9568	1452

Confusion Matrix:

```
[[1370  20]
 [ 38  24]]
```

F1 (Class 1): 0.4528

```
>>> Model: LGBM_Unbal
```

```
Classification Report:
```

	precision	recall	f1-score	support
0.0	0.9724	0.9871	0.9797	1390
1.0	0.5610	0.3710	0.4466	62
accuracy			0.9607	1452
macro avg	0.7667	0.6790	0.7131	1452
weighted avg	0.9548	0.9607	0.9569	1452

```
Confusion Matrix:
```

```
[[1372 18]
```

```
[ 39 23]]
```

```
F1 (Class 1): 0.4466
```

```
💾 Saved best model: RF
```

```
🏆 Best F1 (Class 1): 0.5169
```

```
===== Target: NR-AR-LBD =====
```

```
>>> Model: LR
```

```
Classification Report:
```

	precision	recall	f1-score	support
0.0	0.9854	0.9831	0.9843	1304
1.0	0.5600	0.5957	0.5773	47
accuracy			0.9697	1351
macro avg	0.7727	0.7894	0.7808	1351
weighted avg	0.9706	0.9697	0.9701	1351

```
Confusion Matrix:
```

```
[[1282 22]
```

```
[ 19 28]]
```

```
F1 (Class 1): 0.5773
```

```
>>> Model: KNN
```

```
Classification Report:
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0.0	0.9745	0.9962	0.9852	1304
1.0	0.7222	0.2766	0.4000	47
accuracy			0.9711	1351
macro avg	0.8484	0.6364	0.6926	1351
weighted avg	0.9657	0.9711	0.9649	1351

Confusion Matrix:

```
[[1299  5]
 [ 34 13]]
```

F1 (Class 1): 0.4000

>>> Model: RF

Classification Report:

	precision	recall	f1-score	support
0.0	0.9826	0.9954	0.9890	1304
1.0	0.8000	0.5106	0.6234	47
accuracy			0.9785	1351
macro avg	0.8913	0.7530	0.8062	1351
weighted avg	0.9762	0.9785	0.9762	1351

Confusion Matrix:

```
[[1298  6]
 [ 23 24]]
```

F1 (Class 1): 0.6234

>>> Model: XGB

Classification Report:

	precision	recall	f1-score	support
0.0	0.9855	0.9931	0.9893	1304
1.0	0.7568	0.5957	0.6667	47
accuracy			0.9793	1351
macro avg	0.8711	0.7944	0.8280	1351
weighted avg	0.9776	0.9793	0.9781	1351

Confusion Matrix:

```
[[1295  9]]
```

```
[ 19 28]]  
F1 (Class 1): 0.6667
```

```
>>> Model: LGBM_Bal  
Classification Report:  
precision    recall    f1-score   support  
  
      0.0      0.9870     0.9900     0.9885     1304  
      1.0      0.6977     0.6383     0.6667      47  
  
accuracy          0.9778     1351  
macro avg       0.8423     0.8142     0.8276     1351  
weighted avg     0.9769     0.9778     0.9773     1351
```

Confusion Matrix:

```
[[1291 13]  
 [ 17 30]]  
F1 (Class 1): 0.6667
```

```
>>> Model: LGBM_Unbal  
Classification Report:  
precision    recall    f1-score   support  
  
      0.0      0.9870     0.9885     0.9877     1304  
      1.0      0.6667     0.6383     0.6522      47  
  
accuracy          0.9763     1351  
macro avg       0.8268     0.8134     0.8200     1351  
weighted avg     0.9758     0.9763     0.9761     1351
```

Confusion Matrix:

```
[[1289 15]  
 [ 17 30]]  
F1 (Class 1): 0.6522
```

```
💾 Saved best model: XGB  
🏆 Best F1 (Class 1): 0.6667
```

```
===== Target: NR-AhR =====
```

```
>>> Model: LR
```

Classification Report:

	precision	recall	f1-score	support
0.0	0.9259	0.9307	0.9283	1155
1.0	0.4595	0.4416	0.4503	154
accuracy			0.8732	1309
macro avg	0.6927	0.6861	0.6893	1309
weighted avg	0.8710	0.8732	0.8721	1309

Confusion Matrix:

```
[[1075  80]
 [ 86  68]]
F1 (Class 1): 0.4503
```

>>> Model: KNN

Classification Report:

	precision	recall	f1-score	support
0.0	0.8896	0.9974	0.9404	1155
1.0	0.7857	0.0714	0.1310	154
accuracy			0.8885	1309
macro avg	0.8376	0.5344	0.5357	1309
weighted avg	0.8774	0.8885	0.8452	1309

Confusion Matrix:

```
[[1152   3]
 [ 143  11]]
F1 (Class 1): 0.1310
```

>>> Model: RF

Classification Report:

	precision	recall	f1-score	support
0.0	0.9048	0.9957	0.9481	1155
1.0	0.8684	0.2143	0.3438	154
accuracy			0.9037	1309
macro avg	0.8866	0.6050	0.6459	1309
weighted avg	0.9005	0.9037	0.8770	1309

Confusion Matrix:

```
[[1150    5]
 [ 121   33]]
```

F1 (Class 1): 0.3438

>>> Model: XGB

Classification Report:

	precision	recall	f1-score	support
0.0	0.9206	0.9835	0.9510	1155
1.0	0.7467	0.3636	0.4891	154
accuracy			0.9106	1309
macro avg	0.8336	0.6736	0.7201	1309
weighted avg	0.9001	0.9106	0.8967	1309

Confusion Matrix:

```
[[1136   19]
 [ 98   56]]
```

F1 (Class 1): 0.4891

>>> Model: LGBM_Bal

Classification Report:

	precision	recall	f1-score	support
0.0	0.9503	0.9273	0.9387	1155
1.0	0.5385	0.6364	0.5833	154
accuracy			0.8930	1309
macro avg	0.7444	0.7818	0.7610	1309
weighted avg	0.9019	0.8930	0.8968	1309

Confusion Matrix:

```
[[1071   84]
 [ 56   98]]
```

F1 (Class 1): 0.5833

>>> Model: LGBM_Unbal

Classification Report:

precision	recall	f1-score	support
-----------	--------	----------	---------

0.0	0.9486	0.9264	0.9374	1155
1.0	0.5304	0.6234	0.5731	154
accuracy			0.8908	1309
macro avg	0.7395	0.7749	0.7552	1309
weighted avg	0.8994	0.8908	0.8945	1309

Confusion Matrix:

```
[[1070  85]
 [ 58  96]]
F1 (Class 1): 0.5731
```

```
💾 Saved best model: LGBM_Bal
🏆 Best F1 (Class 1): 0.5833
```

===== Target: NR-Aromatase =====

>>> Model: LR

Classification Report:

	precision	recall	f1-score	support
0.0	0.9610	0.9837	0.9722	1103
1.0	0.4706	0.2667	0.3404	60
accuracy			0.9467	1163
macro avg	0.7158	0.6252	0.6563	1163
weighted avg	0.9357	0.9467	0.9396	1163

Confusion Matrix:

```
[[1085  18]
 [ 44  16]]
F1 (Class 1): 0.3404
```

>>> Model: KNN

Classification Report:

	precision	recall	f1-score	support
0.0	0.9533	0.9991	0.9757	1103
1.0	0.8571	0.1000	0.1791	60

accuracy			0.9527	1163
macro avg	0.9052	0.5495	0.5774	1163
weighted avg	0.9483	0.9527	0.9346	1163

Confusion Matrix:

```
[[1102    1]
 [ 54    6]]
```

F1 (Class 1): 0.1791

>>> Model: RF

Classification Report:

	precision	recall	f1-score	support
0.0	0.9565	0.9973	0.9765	1103
1.0	0.7692	0.1667	0.2740	60
accuracy			0.9544	1163
macro avg	0.8629	0.5820	0.6252	1163
weighted avg	0.9469	0.9544	0.9402	1163

Confusion Matrix:

```
[[1100    3]
 [ 50   10]]
```

F1 (Class 1): 0.2740

>>> Model: XGB

Classification Report:

	precision	recall	f1-score	support
0.0	0.9588	0.9918	0.9750	1103
1.0	0.5909	0.2167	0.3171	60
accuracy			0.9518	1163
macro avg	0.7749	0.6043	0.6461	1163
weighted avg	0.9398	0.9518	0.9411	1163

Confusion Matrix:

```
[[1094    9]
 [ 47   13]]
```

F1 (Class 1): 0.3171

>>> Model: LGBM_Bal

Classification Report:

	precision	recall	f1-score	support
0.0	0.9641	0.9746	0.9693	1103
1.0	0.4167	0.3333	0.3704	60
accuracy			0.9415	1163
macro avg	0.6904	0.6540	0.6699	1163
weighted avg	0.9359	0.9415	0.9384	1163

Confusion Matrix:

```
[[1075  28]
 [ 40  20]]
```

F1 (Class 1): 0.3704

>>> Model: LGBM_Unbal

Classification Report:

	precision	recall	f1-score	support
0.0	0.9644	0.9837	0.9740	1103
1.0	0.5263	0.3333	0.4082	60
accuracy			0.9501	1163
macro avg	0.7454	0.6585	0.6911	1163
weighted avg	0.9418	0.9501	0.9448	1163

Confusion Matrix:

```
[[1085  18]
 [ 40  20]]
```

F1 (Class 1): 0.4082

Saved best model: LGBM_Unbal

Best F1 (Class 1): 0.4082

===== Target: NR-ER =====

>>> Model: LR

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0.0	0.9058	0.8722	0.8887	1080
1.0	0.3030	0.3797	0.3371	158
accuracy			0.8094	1238
macro avg	0.6044	0.6260	0.6129	1238
weighted avg	0.8288	0.8094	0.8183	1238

Confusion Matrix:

```
[[942 138]
 [ 98  60]]
```

F1 (Class 1): 0.3371

>>> Model: KNN

Classification Report:

	precision	recall	f1-score	support
0.0	0.8869	0.9944	0.9376	1080
1.0	0.7778	0.1329	0.2270	158
accuracy			0.8845	1238
macro avg	0.8323	0.5637	0.5823	1238
weighted avg	0.8729	0.8845	0.8469	1238

Confusion Matrix:

```
[[1074   6]
 [ 137  21]]
```

F1 (Class 1): 0.2270

>>> Model: RF

Classification Report:

	precision	recall	f1-score	support
0.0	0.8944	0.9963	0.9426	1080
1.0	0.8857	0.1962	0.3212	158
accuracy			0.8942	1238
macro avg	0.8901	0.5962	0.6319	1238
weighted avg	0.8933	0.8942	0.8633	1238

Confusion Matrix:

```
[[1076   4]]
```

```
[ 127  31]]  
F1 (Class 1): 0.3212
```

```
>>> Model: XGB  
Classification Report:  
precision    recall   f1-score   support  
  
      0.0      0.9048    0.9861    0.9437     1080  
      1.0      0.7541    0.2911    0.4201      158  
  
accuracy          0.8974  
macro avg      0.8295    0.6386    0.6819     1238  
weighted avg    0.8856    0.8974    0.8769     1238
```

Confusion Matrix:

```
[[1065  15]  
 [ 112  46]]  
F1 (Class 1): 0.4201
```

```
>>> Model: LGBM_Bal  
Classification Report:  
precision    recall   f1-score   support  
  
      0.0      0.9175    0.9065    0.9120     1080  
      1.0      0.4094    0.4430    0.4255      158  
  
accuracy          0.8473  
macro avg      0.6634    0.6748    0.6688     1238  
weighted avg    0.8527    0.8473    0.8499     1238
```

Confusion Matrix:

```
[[979 101]  
 [ 88  70]]  
F1 (Class 1): 0.4255
```

```
>>> Model: LGBM_Unbal  
Classification Report:  
precision    recall   f1-score   support  
  
      0.0      0.9173    0.9037    0.9104     1080  
      1.0      0.4023    0.4430    0.4217      158
```

accuracy		0.8449	1238
macro avg	0.6598	0.6734	0.6661 1238
weighted avg	0.8516	0.8449	0.8481 1238

Confusion Matrix:

```
[[976 104]
 [ 88  70]]
```

F1 (Class 1): 0.4217

💾 Saved best model: LGBM_Bal
🏆 Best F1 (Class 1): 0.4255

===== Target: NR-ER-LBD =====

>>> Model: LR

Classification Report:

	precision	recall	f1-score	support
0.0	0.9651	0.9833	0.9741	1320
1.0	0.5111	0.3286	0.4000	70
accuracy			0.9504	1390
macro avg	0.7381	0.6560	0.6871	1390
weighted avg	0.9422	0.9504	0.9452	1390

Confusion Matrix:

```
[[1298  22]
 [ 47  23]]
```

F1 (Class 1): 0.4000

>>> Model: KNN

Classification Report:

	precision	recall	f1-score	support
0.0	0.9550	0.9970	0.9755	1320
1.0	0.6667	0.1143	0.1951	70
accuracy			0.9525	1390
macro avg	0.8108	0.5556	0.5853	1390
weighted avg	0.9405	0.9525	0.9362	1390

Confusion Matrix:

```
[[1316    4]
 [ 62     8]]
```

F1 (Class 1): 0.1951

>>> Model: RF

Classification Report:

	precision	recall	f1-score	support
0.0	0.9620	0.9970	0.9792	1320
1.0	0.8182	0.2571	0.3913	70
accuracy			0.9597	1390
macro avg	0.8901	0.6271	0.6852	1390
weighted avg	0.9547	0.9597	0.9496	1390

Confusion Matrix:

```
[[1316    4]
 [ 52    18]]
```

F1 (Class 1): 0.3913

>>> Model: XGB

Classification Report:

	precision	recall	f1-score	support
0.0	0.9661	0.9932	0.9795	1320
1.0	0.7273	0.3429	0.4660	70
accuracy			0.9604	1390
macro avg	0.8467	0.6680	0.7227	1390
weighted avg	0.9541	0.9604	0.9536	1390

Confusion Matrix:

```
[[1311    9]
 [ 46    24]]
```

F1 (Class 1): 0.4660

>>> Model: LGBM_Bal

Classification Report:

precision	recall	f1-score	support
-----------	--------	----------	---------

0.0	0.9727	0.9705	0.9716	1320
1.0	0.4658	0.4857	0.4755	70
accuracy			0.9460	1390
macro avg	0.7192	0.7281	0.7235	1390
weighted avg	0.9471	0.9460	0.9466	1390

Confusion Matrix:

```
[[1281  39]
 [ 36  34]]
```

F1 (Class 1): 0.4755

>>> Model: LGBM_Unbal

Classification Report:

	precision	recall	f1-score	support
0.0	0.9712	0.9720	0.9716	1320
1.0	0.4638	0.4571	0.4604	70
accuracy			0.9460	1390
macro avg	0.7175	0.7146	0.7160	1390
weighted avg	0.9457	0.9460	0.9459	1390

Confusion Matrix:

```
[[1283  37]
 [ 38  32]]
```

F1 (Class 1): 0.4604

💾 Saved best model: LGBM_Bal
🏆 Best F1 (Class 1): 0.4755

===== Target: NR-PPAR-gamma =====

>>> Model: LR

Classification Report:

	precision	recall	f1-score	support
0.0	0.9756	0.9920	0.9838	1252
1.0	0.3750	0.1622	0.2264	37

accuracy			0.9682	1289
macro avg	0.6753	0.5771	0.6051	1289
weighted avg	0.9584	0.9682	0.9620	1289

Confusion Matrix:

```
[[1242 10]
 [ 31  6]]
```

F1 (Class 1): 0.2264

>>> Model: KNN

Classification Report:

	precision	recall	f1-score	support
0.0	0.9728	1.0000	0.9862	1252
1.0	1.0000	0.0541	0.1026	37
accuracy			0.9728	1289
macro avg	0.9864	0.5270	0.5444	1289
weighted avg	0.9736	0.9728	0.9609	1289

Confusion Matrix:

```
[[1252 0]
 [ 35 2]]
```

F1 (Class 1): 0.1026

>>> Model: RF

Classification Report:

	precision	recall	f1-score	support
0.0	0.9728	1.0000	0.9862	1252
1.0	1.0000	0.0541	0.1026	37
accuracy			0.9728	1289
macro avg	0.9864	0.5270	0.5444	1289
weighted avg	0.9736	0.9728	0.9609	1289

Confusion Matrix:

```
[[1252 0]
 [ 35 2]]
```

F1 (Class 1): 0.1026

>>> Model: XGB

Classification Report:

	precision	recall	f1-score	support
0.0	0.9757	0.9952	0.9854	1252
1.0	0.5000	0.1622	0.2449	37
accuracy			0.9713	1289
macro avg	0.7379	0.5787	0.6151	1289
weighted avg	0.9621	0.9713	0.9641	1289

Confusion Matrix:

```
[[1246  6]
 [ 31  6]]
```

F1 (Class 1): 0.2449

>>> Model: LGBM_Bal

Classification Report:

	precision	recall	f1-score	support
0.0	0.9794	0.9896	0.9845	1252
1.0	0.4583	0.2973	0.3607	37
accuracy			0.9697	1289
macro avg	0.7189	0.6435	0.6726	1289
weighted avg	0.9645	0.9697	0.9666	1289

Confusion Matrix:

```
[[1239 13]
 [ 26 11]]
```

F1 (Class 1): 0.3607

>>> Model: LGBM_Unbal

Classification Report:

	precision	recall	f1-score	support
0.0	0.9794	0.9888	0.9841	1252
1.0	0.4400	0.2973	0.3548	37
accuracy			0.9690	1289
macro avg	0.7097	0.6431	0.6695	1289

weighted avg	0.9639	0.9690	0.9660	1289
--------------	--------	--------	--------	------

Confusion Matrix:

```
[[1238 14]
 [ 26 11]]
```

F1 (Class 1): 0.3548

Saved best model: LGBM_Bal

Best F1 (Class 1): 0.3607

===== Target: SR-ARE =====

>>> Model: LR

Classification Report:

	precision	recall	f1-score	support
0.0	0.8839	0.8495	0.8664	977
1.0	0.3496	0.4202	0.3816	188
accuracy			0.7803	1165
macro avg	0.6167	0.6349	0.6240	1165
weighted avg	0.7977	0.7803	0.7882	1165

Confusion Matrix:

```
[[830 147]
 [109 79]]
```

F1 (Class 1): 0.3816

>>> Model: KNN

Classification Report:

	precision	recall	f1-score	support
0.0	0.8434	0.9980	0.9142	977
1.0	0.7778	0.0372	0.0711	188
accuracy			0.8429	1165
macro avg	0.8106	0.5176	0.4926	1165
weighted avg	0.8328	0.8429	0.7781	1165

Confusion Matrix:

```
[[975 2]]
```

```
[181 7]  
F1 (Class 1): 0.0711
```

```
>>> Model: RF  
Classification Report:  
precision recall f1-score support  
  
0.0 0.8604 0.9969 0.9237 977  
1.0 0.9091 0.1596 0.2715 188  
  
accuracy 0.8618 1165  
macro avg 0.8848 0.5783 0.5976 1165  
weighted avg 0.8683 0.8618 0.8184 1165
```

Confusion Matrix:

```
[[974 3]  
 [158 30]]  
F1 (Class 1): 0.2715
```

```
>>> Model: XGB  
Classification Report:  
precision recall f1-score support  
  
0.0 0.8890 0.9754 0.9302 977  
1.0 0.7419 0.3670 0.4911 188  
  
accuracy 0.8773 1165  
macro avg 0.8155 0.6712 0.7107 1165  
weighted avg 0.8653 0.8773 0.8593 1165
```

Confusion Matrix:

```
[[953 24]  
 [119 69]]  
F1 (Class 1): 0.4911
```

```
>>> Model: LGBM_Bal  
Classification Report:  
precision recall f1-score support  
  
0.0 0.9140 0.8813 0.8973 977  
1.0 0.4798 0.5691 0.5207 188
```

accuracy		0.8309	1165
macro avg	0.6969	0.7252	0.7090
weighted avg	0.8439	0.8309	0.8366

Confusion Matrix:

```
[[861 116]
 [ 81 107]]
```

F1 (Class 1): 0.5207

>>> Model: LGBM_Unbal

Classification Report:

	precision	recall	f1-score	support
0.0	0.9151	0.8710	0.8925	977
1.0	0.4638	0.5798	0.5154	188
accuracy			0.8240	1165
macro avg	0.6894	0.7254	0.7039	1165
weighted avg	0.8422	0.8240	0.8316	1165

Confusion Matrix:

```
[[851 126]
 [ 79 109]]
```

F1 (Class 1): 0.5154

💾 Saved best model: LGBM_Bal

🏆 Best F1 (Class 1): 0.5207

===== Target: SR-ATAD5 =====

>>> Model: LR

Classification Report:

	precision	recall	f1-score	support
0.0	0.9682	0.9846	0.9763	1360
1.0	0.3000	0.1698	0.2169	53
accuracy			0.9540	1413
macro avg	0.6341	0.5772	0.5966	1413
weighted avg	0.9431	0.9540	0.9478	1413

Confusion Matrix:

```
[[1339  21]
 [ 44   9]]
```

F1 (Class 1): 0.2169

>>> Model: KNN

Classification Report:

	precision	recall	f1-score	support
0.0	0.9658	0.9978	0.9816	1360
1.0	0.6250	0.0943	0.1639	53
accuracy			0.9639	1413
macro avg	0.7954	0.5461	0.5727	1413
weighted avg	0.9531	0.9639	0.9509	1413

Confusion Matrix:

```
[[1357  3]
 [ 48   5]]
```

F1 (Class 1): 0.1639

>>> Model: RF

Classification Report:

	precision	recall	f1-score	support
0.0	0.9652	1.0000	0.9823	1360
1.0	1.0000	0.0755	0.1404	53
accuracy			0.9653	1413
macro avg	0.9826	0.5377	0.5613	1413
weighted avg	0.9665	0.9653	0.9507	1413

Confusion Matrix:

```
[[1360  0]
 [ 49   4]]
```

F1 (Class 1): 0.1404

>>> Model: XGB

Classification Report:

precision	recall	f1-score	support
-----------	--------	----------	---------

0.0	0.9664	0.9949	0.9804	1360
1.0	0.4615	0.1132	0.1818	53
accuracy			0.9618	1413
macro avg	0.7140	0.5540	0.5811	1413
weighted avg	0.9475	0.9618	0.9505	1413

Confusion Matrix:

```
[[1353  7]
 [ 47  6]]
```

F1 (Class 1): 0.1818

>>> Model: LGBM_Bal

Classification Report:

	precision	recall	f1-score	support
0.0	0.9701	0.9787	0.9744	1360
1.0	0.2927	0.2264	0.2553	53
accuracy			0.9505	1413
macro avg	0.6314	0.6025	0.6148	1413
weighted avg	0.9447	0.9505	0.9474	1413

Confusion Matrix:

```
[[1331 29]
 [ 41 12]]
```

F1 (Class 1): 0.2553

>>> Model: LGBM_Unbal

Classification Report:

	precision	recall	f1-score	support
0.0	0.9688	0.9824	0.9755	1360
1.0	0.2941	0.1887	0.2299	53
accuracy			0.9526	1413
macro avg	0.6315	0.5855	0.6027	1413
weighted avg	0.9435	0.9526	0.9476	1413

Confusion Matrix:

```
[[1336 24]
 [ 43 10]]
F1 (Class 1): 0.2299
```

```
💾 Saved best model: LGBM_Bal
🏆 Best F1 (Class 1): 0.2553
```

```
===== Target: SR-HSE =====
```

```
>>> Model: LR
Classification Report:
      precision    recall   f1-score   support
          0.0       0.9528   0.9614   0.9571     1218
          1.0       0.2540   0.2162   0.2336      74
accuracy                           0.9187     1292
macro avg       0.6034   0.5888   0.5953     1292
weighted avg     0.9128   0.9187   0.9157     1292
```

Confusion Matrix:

```
[[1171 47]
 [ 58 16]]
F1 (Class 1): 0.2336
```

```
>>> Model: KNN
Classification Report:
      precision    recall   f1-score   support
          0.0       0.9492   0.9975   0.9728     1218
          1.0       0.7500   0.1216   0.2093      74
accuracy                           0.9474     1292
macro avg       0.8496   0.5596   0.5910     1292
weighted avg     0.9378   0.9474   0.9290     1292
```

Confusion Matrix:

```
[[1215 3]
 [ 65 9]]
F1 (Class 1): 0.2093
```

>>> Model: RF

Classification Report:

	precision	recall	f1-score	support
0.0	0.9478	0.9992	0.9728	1218
1.0	0.8750	0.0946	0.1707	74
accuracy			0.9474	1292
macro avg	0.9114	0.5469	0.5718	1292
weighted avg	0.9436	0.9474	0.9269	1292

Confusion Matrix:

```
[[1217  1]
 [ 67  7]]
F1 (Class 1): 0.1707
```

>>> Model: XGB

Classification Report:

	precision	recall	f1-score	support
0.0	0.9543	0.9934	0.9735	1218
1.0	0.6667	0.2162	0.3265	74
accuracy			0.9489	1292
macro avg	0.8105	0.6048	0.6500	1292
weighted avg	0.9378	0.9489	0.9364	1292

Confusion Matrix:

```
[[1210  8]
 [ 58 16]]
F1 (Class 1): 0.3265
```

>>> Model: LGBM_Bal

Classification Report:

	precision	recall	f1-score	support
0.0	0.9625	0.9696	0.9661	1218
1.0	0.4308	0.3784	0.4029	74
accuracy			0.9358	1292
macro avg	0.6966	0.6740	0.6845	1292

```
weighted avg    0.9321    0.9358    0.9338    1292
```

Confusion Matrix:

```
[[1181  37]
 [ 46  28]]
```

F1 (Class 1): 0.4029

>>> Model: LGBM_Unbal

Classification Report:

	precision	recall	f1-score	support
0.0	0.9600	0.9655	0.9628	1218
1.0	0.3731	0.3378	0.3546	74
accuracy			0.9296	1292
macro avg	0.6666	0.6517	0.6587	1292
weighted avg	0.9264	0.9296	0.9279	1292

Confusion Matrix:

```
[[1176  42]
 [ 49  25]]
```

F1 (Class 1): 0.3546

Saved best model: LGBM_Bal

Best F1 (Class 1): 0.4029

===== Target: SR-MMP =====

>>> Model: LR

Classification Report:

	precision	recall	f1-score	support
0.0	0.9114	0.9263	0.9188	977
1.0	0.5714	0.5217	0.5455	184
accuracy			0.8622	1161
macro avg	0.7414	0.7240	0.7321	1161
weighted avg	0.8575	0.8622	0.8596	1161

Confusion Matrix:

```
[[905  72]
```

```
[ 88  96]  
F1 (Class 1): 0.5455
```

```
>>> Model: KNN  
Classification Report:  
precision    recall   f1-score   support  
  
      0.0      0.8526    0.9949    0.9183      977  
      1.0      0.7619    0.0870    0.1561     184  
  
accuracy          0.8510      1161  
macro avg       0.8073    0.5409    0.5372     1161  
weighted avg     0.8383    0.8510    0.7975     1161
```

Confusion Matrix:

```
[[972  5]  
 [168 16]]  
F1 (Class 1): 0.1561
```

```
>>> Model: RF  
Classification Report:  
precision    recall   f1-score   support  
  
      0.0      0.8862    0.9959    0.9378      977  
      1.0      0.9365    0.3207    0.4777     184  
  
accuracy          0.8889      1161  
macro avg       0.9113    0.6583    0.7078     1161  
weighted avg     0.8941    0.8889    0.8649     1161
```

Confusion Matrix:

```
[[973  4]  
 [125 59]]  
F1 (Class 1): 0.4777
```

```
>>> Model: XGB  
Classification Report:  
precision    recall   f1-score   support  
  
      0.0      0.9209    0.9652    0.9425      977  
      1.0      0.7518    0.5598    0.6417     184
```

accuracy		0.9009	1161
macro avg	0.8364	0.7625	0.7921
weighted avg	0.8941	0.9009	0.8949
			1161

Confusion Matrix:

```
[[943 34]
 [ 81 103]]
```

F1 (Class 1): 0.6417

>>> Model: LGBM_Bal

Classification Report:

	precision	recall	f1-score	support
0.0	0.9411	0.9161	0.9284	977
1.0	0.6095	0.6957	0.6497	184
accuracy			0.8811	1161
macro avg	0.7753	0.8059	0.7891	1161
weighted avg	0.8886	0.8811	0.8843	1161

Confusion Matrix:

```
[[895 82]
 [ 56 128]]
```

F1 (Class 1): 0.6497

>>> Model: LGBM_Unbal

Classification Report:

	precision	recall	f1-score	support
0.0	0.9436	0.9243	0.9338	977
1.0	0.6373	0.7065	0.6701	184
accuracy			0.8898	1161
macro avg	0.7904	0.8154	0.8020	1161
weighted avg	0.8950	0.8898	0.8920	1161

Confusion Matrix:

```
[[903 74]
 [ 54 130]]
```

F1 (Class 1): 0.6701

💾 Saved best model: LGBM_Unbal
🏆 Best F1 (Class 1): 0.6701

===== Target: SR-p53 =====

>>> Model: LR

Classification Report:

	precision	recall	f1-score	support
0.0	0.9529	0.9574	0.9552	1269
1.0	0.3165	0.2941	0.3049	85
accuracy			0.9158	1354
macro avg	0.6347	0.6258	0.6300	1354
weighted avg	0.9130	0.9158	0.9144	1354

Confusion Matrix:

[[1215 54]
 [60 25]]

F1 (Class 1): 0.3049

>>> Model: KNN

Classification Report:

	precision	recall	f1-score	support
0.0	0.9393	0.9992	0.9683	1269
1.0	0.7500	0.0353	0.0674	85
accuracy			0.9387	1354
macro avg	0.8446	0.5173	0.5179	1354
weighted avg	0.9274	0.9387	0.9118	1354

Confusion Matrix:

[[1268 1]
 [82 3]]

F1 (Class 1): 0.0674

>>> Model: RF

Classification Report:

precision recall f1-score support

0.0	0.9441	0.9976	0.9701	1269
1.0	0.7692	0.1176	0.2041	85
accuracy			0.9424	1354
macro avg	0.8567	0.5576	0.5871	1354
weighted avg	0.9331	0.9424	0.9220	1354

Confusion Matrix:

```
[[1266  3]
 [ 75 10]]
F1 (Class 1): 0.2041
```

>>> Model: XGB

Classification Report:

	precision	recall	f1-score	support
0.0	0.9500	0.9882	0.9687	1269
1.0	0.5588	0.2235	0.3193	85
accuracy			0.9402	1354
macro avg	0.7544	0.6059	0.6440	1354
weighted avg	0.9254	0.9402	0.9279	1354

Confusion Matrix:

```
[[1254 15]
 [ 66 19]]
F1 (Class 1): 0.3193
```

>>> Model: LGBM_Bal

Classification Report:

	precision	recall	f1-score	support
0.0	0.9612	0.9574	0.9593	1269
1.0	0.4000	0.4235	0.4114	85
accuracy			0.9239	1354
macro avg	0.6806	0.6905	0.6854	1354
weighted avg	0.9260	0.9239	0.9249	1354

Confusion Matrix:

```
[[1215  54]
 [ 49  36]]
F1 (Class 1): 0.4114
```

```
>>> Model: LGBM_Unbal
Classification Report:
precision    recall    f1-score   support

      0.0      0.9651    0.9582    0.9616     1269
      1.0      0.4362    0.4824    0.4581      85

  accuracy                           0.9284    1354
 macro avg      0.7006    0.7203    0.7099    1354
weighted avg     0.9319    0.9284    0.9300    1354
```

Confusion Matrix:

```
[[1216  53]
 [ 44  41]]
F1 (Class 1): 0.4581
```

```
💾 Saved best model: LGBM_Unbal
🏆 Best F1 (Class 1): 0.4581
```

In [130...]

```
results_df_normalize = pd.DataFrame(results).T

numeric_cols = results_df_normalize.select_dtypes(include="number")

results_df_normalize["Best_Model"] = numeric_cols.idxmax(axis=1)
results_df_normalize["Best_CV_F1"] = numeric_cols.max(axis=1)

results_df_normalize.to_csv("results_with_normalize.csv", index=False)
display(results_df_normalize.sort_values("Best_CV_F1", ascending=False))
```

	LR	KNN	RF	XGB	LGBM_Bal	LGBM_Unbal	Best_Model	Best_CV_F1
SR-MMP	0.545455	0.156098	0.477733	0.641745	0.649746	0.670103	LGBM_Unbal	0.670103
NR-AR-LBD	0.577320	0.400000	0.623377	0.666667	0.666667	0.652174	XGB	0.666667
NR-AhR	0.450331	0.130952	0.343750	0.489083	0.583333	0.573134	LGBM_Bal	0.583333
SR-ARE	0.381643	0.071066	0.271493	0.491103	0.520681	0.515366	LGBM_Bal	0.520681
NR-AR	0.389381	0.430380	0.516854	0.500000	0.452830	0.446602	RF	0.516854
NR-ER-LBD	0.400000	0.195122	0.391304	0.466019	0.475524	0.460432	LGBM_Bal	0.475524
SR-p53	0.304878	0.067416	0.204082	0.319328	0.411429	0.458101	LGBM_Unbal	0.458101
NR-ER	0.337079	0.227027	0.321244	0.420091	0.425532	0.421687	LGBM_Bal	0.425532
NR-Aromatase	0.340426	0.179104	0.273973	0.317073	0.370370	0.408163	LGBM_Unbal	0.408163
SR-HSE	0.233577	0.209302	0.170732	0.326531	0.402878	0.354610	LGBM_Bal	0.402878
NR-PPAR-gamma	0.226415	0.102564	0.102564	0.244898	0.360656	0.354839	LGBM_Bal	0.360656
SR-ATAD5	0.216867	0.163934	0.140351	0.181818	0.255319	0.229885	LGBM_Bal	0.255319

Why Power Transformation?

- Molecular descriptors are naturally right-skewed
- Linear models (LR) and distance-based models (KNN) assume normally distributed features
- Yeo-Johnson handles zero values (unlike Box-Cox)
- **Critical:** Fitted only on training data to prevent leakage

🔍 Key Observations

- Each toxicity assay was treated as an independent binary classification problem.
- Multiple machine learning models were evaluated, including Logistic Regression, Random Forest, KNN, XGBoost, and LightGBM.
- Stratified K-Fold cross-validation was used to handle class imbalance during model evaluation.
- F1 score was selected as the primary metric to balance precision and recall for the minority (toxic) class.

- LightGBM (balanced or unbalanced) and Random Forest models consistently showed the best performance across most assays.
 - For each assay, the model with the highest cross-validated F1 score was selected as the final best model.
 - Normalized molecular descriptors generally improved or maintained model performance compared to raw features.
 - The selected best models were retrained on the full available dataset for each assay.
 - Final trained models were saved to disk for future prediction, reuse, or deployment workflows.
-

Limitations & Future Work

- **Missing data:** 7-26% of labels missing per assay (handled by per-target modeling)
- **Class imbalance:** Minority class as low as 2.88% (addressed via F1 metric and balanced sampling)
- **External validation:** Models not tested on external datasets
- **Feature engineering:** Could explore 3D descriptors, graph neural networks, or transformer-based molecular embeddings

Future improvements:

- Multi-task learning to leverage correlations between assays
 - Ensemble stacking of best models
-

Project Conclusion

This project successfully built machine learning models to predict chemical toxicity across 12 different biological tests using molecular structures.

What we did:

- Converted 7,823 chemical formulas (SMILES) into numerical features that computers can understand
- Created two types of features: basic properties (like weight, size) and structural fingerprints (molecular patterns)
- Applied power transformation to make data more normally distributed, which helps models learn better
- Tested 6 different machine learning algorithms on each of the 12 toxicity tests

- Selected the best-performing model for each test based on F1 score (balances precision and recall)

Key findings:

- **LightGBM and Random Forest** worked best overall, handling the imbalanced data effectively
- **F1 scores ranged from 0.26 to 0.67**, with SR-MMP toxicity being easiest to predict (67% F1)
- **Power transformation improved results** for models like Logistic Regression and KNN
- **Proper data handling mattered** - we transformed features ONLY on training data to avoid leakage

Challenges addressed:

- **Missing labels:** 7-26% of molecules weren't tested for every assay - solved by building separate models for each target
- **Class imbalance:** Only 3-16% of molecules were toxic - handled using balanced sampling and F1 scoring
- **Feature scaling:** Different molecular properties had very different ranges - normalized them properly

Real-world impact: These models can now screen new chemicals for toxicity **before** expensive lab testing, speeding up drug discovery and chemical safety assessment. All models are saved and ready for deployment in production environments.

Future improvements:

- Test on external datasets to verify generalization
- Try advanced techniques like graph neural networks that directly learn from molecular structures
- Use multi-task learning to leverage relationships between different toxicity tests